

Research prototype · seeking pre-seed

The compiler that **learns.**

An AI that learns to optimize machine code — by understanding how data actually flows through it, not by reading text. Built on JEPA, a self-supervised world-model approach. To our knowledge, the first applied to assembly.

THE STAKES

The hardest code on earth runs on the thinnest margins.

Trading engines. Physics simulations. Databases. There, **every cycle of speed and every byte of memory matters** — and one thing is never negotiable: the code must be **correct, right down to the metal**.

ns

latency that moves money

100%

correctness, non-negotiable

\$B

spent chasing performance

THE HIDDEN TRUTH

**Performance isn't decided
in your source code.
It's decided in the **assembly**
the compiler produces.**

So why not just let the compiler handle it? That's exactly the problem.

THE PROBLEM · TODAY'S COMPILERS

Frozen heuristics that never learn.

- **Hand-tuned rules** an expert wrote years ago, applied in a **fixed order**.
- **Crude cost models** that approximate the hardware — and drift further from reality with every new chip generation.
- They must guarantee correctness, so they **play it safe** and leave performance on the table.
- They **don't learn**. Every new architecture = more manual tuning, by hand, by experts. The compiler never gets smarter on its own.

OUR INSIGHT

Represent code differently. Learn from it differently.

The structure, not the text. We turn each block of assembly into a **graph of data flow** — what depends on what. That is the real structure that drives performance. Unlike an LLM, we never act on text.

A bet on how it learns. On top of that graph we use **JEPA** — Yann LeCun's self-supervised world-model approach, just beginning to be explored in new domains (e.g. biology). **To our knowledge, no one has applied it to assembly. We're the first.**

HOW IT WORKS

From instructions to a learned world model of code.



The model splits its understanding in two: **what the code does** (z_{sem}) and **its optimization profile** (z_{speed}) — learned with zero manual labels.

PROOF · IT WORKS

The model disentangles meaning from speed.


z_{sem} = what the code does **gap 0.89**

...ignores optimization level **silhouette ≈ 0 ✓**

z_{speed} = optimization profile **isolated axis**

...ignores the program **silhouette -0.91 ✓**

Self-supervised, no labels · ~8k ExeBench functions · ProgramML graphs ·
trained from scratch on one B200 GPU.



z_{sem} clusters by program; z_{speed} by optimization. The two are decorrelated.

PROOF · THE TEAM DOES REAL SCIENCE

We measure honestly, and we debug deep.

- **We gate before we train.** A probe honestly quantified where today's compiler saturates ($O0 \rightarrow O1$ changes 100% of graphs; $O2 \approx O3$ are identical). We report the ceiling instead of hiding it.
- **We found a subtle collapse bug** in our own loss via an **adversarial multi-agent review** — and fixed it.

3

latent dims used (bug)

72

after the fix · same data

Effective rank of z_{sem} , $3 \rightarrow 72$ of 96. The kind of rigor that machine-level correctness demands.

GO-TO-MARKET · THE "GITGUARDIAN" PLAYBOOK

We don't sell promises. We send proof.



Immediate, irrefutable proof of value — converting open source into qualified inbound, the same viral motion that built GitGuardian.

THE VISION

A universal, AI-driven compiler.

The encoder is the foundational brick. The end state: take **any source, in any language**, and compile it **optimally for any hardware** — CPU, GPU, specialized chips. A world model that gets smarter with every program it sees.

WHY NOW

The window just opened.

JEPA

self-supervised world models are crossing into new sciences — untouched in assembly

chips

each new architecture makes hand-tuning compilers more unsustainable

compute

distributed GPU markets make training economical for a small team

EXECUTION · HOW WE RUN LEAN

Massive compute, fraction of the cost.

- **Infrastructure:** distributed cloud (Vast.ai) for massive GPU power at a fraction of hyperscaler cost. Today: a single B200 trains the encoder end-to-end.
- **Data & curriculum:** training structured in **educational echelons** — the model learns algorithmics from the basics up to the most complex architectures.

ProgramML graphs

ExeBench corpus

JEPA / VICReg

PyTorch Geometric

Vast.ai · B200

Full pipeline (probe → cache → train → eval) reproduces from one command line.

WHAT YOU'RE PROBABLY ASKING

The hard questions, answered honestly.

- "Does this actually make code faster yet?"

Not yet — we've proven the model **separates meaning from optimization**. Next milestone: first measured speedup on a real benchmark.

- "The optimization signal looks tiny."

On isolated functions, yes — we measured it. **Whole programs** (inlining, vectorization) are where the gains live; that's the corpus we scale to.

- "Why won't LLVM / Google / an LLM lab crush you?"

We learn on the **data-flow graph, not text**. A JEPA world model + proprietary corpus + the optimizer loop is the moat — not a weekend clone.

- "How do you guarantee correctness?"

We **propose, the toolchain proves**: every AI rewrite is gated by the compiler's equivalence checks. We never ship an unverified transform.

WHAT YOUR INVESTMENT DE-RISKS

Three milestones from representation to speedup.

1

First beat-O3 win.

One closed loop — encoder → search → codegen
→ verify → measured wall-clock — that beats
clang -O3 on a real hot loop, correctness-checked.

2

Headroom proof.

Re-run the gate on **whole-program** corpora (SPEC
/ cBench) to show the exploitable runtime delta
exists where the thesis lives. Days, not months.

3

Safety rail + moat.

Propose-then-verify with zero correctness
regressions, and the first 10k (**graph** → **measured
speedup**) records — the data moat incumbents
don't have.

| We know exactly what's unproven. This is the plan to prove it.

Let's make the compiler learn.

Seeking pre-seed to take the encoder from **proven disentanglement** to a
first measured speedup on real code.

[◀ contact](#) · [team](#) · [the ask](#) ▶